

```
In [37]: import numpy as np
import pandas as pd

from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split, GridSearchCV, RandomizedSearchCV
from sklearn.metrics import confusion_matrix, accuracy_score, classification_report
from sklearn.tree import DecisionTreeClassifier, plot_tree
from sklearn.ensemble import RandomForestClassifier, AdaBoostClassifier
import seaborn as sns
import matplotlib.pyplot as plt
```

```
In [2]: #import numpy as np
#import pandas as pd
#from sklearn.linear_model import LogisticRegression
#from sklearn.model_selection import train_test_split, GridSearchCV, RandomizedSearchCV
#from sklearn.metrics import confusion_matrix, accuracy_score, classification_report
#from sklearn.tree import DecisionTreeClassifier, plot_tree
#from sklearn.ensemble import RandomForestClassifier, AdaBoostClassifier
#import matplotlib.pyplot as plt
#import seaborn as sns
```

## 1: Problem Statement

```
In [3]: #Target Column="insuranceclaim"(Statistical Binary Classification)
```

## 2: Data Gathering

```
In [4]: df=pd.read_csv(r"C:\Users\Dell\Downloads\insurance2.csv")
df
```

```
Out[4]:
```

	age	sex	bmi	children	smoker	region	charges	insuranceclaim
0	19	0	27.900	0	1	3	16884.92400	1
1	18	1	33.770	1	0	2	1725.55230	1
2	28	1	33.000	3	0	2	4449.46200	0
3	33	1	22.705	0	0	1	21984.47061	0
4	32	1	28.880	0	0	1	3866.85520	1
...	...	...	...	...	...	...	...	...
1333	50	1	30.970	3	0	1	10600.54830	0
1334	18	0	31.920	0	0	0	2205.98080	1
1335	18	0	36.850	0	0	2	1629.83350	1
1336	21	0	25.800	0	0	3	2007.94500	0
1337	61	0	29.070	0	1	1	29141.36030	1

1338 rows × 8 columns

### 3: EDA(Exploratory data analysis)

```
In [5]: df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1338 entries, 0 to 1337
Data columns (total 8 columns):
 #   Column          Non-Null Count  Dtype  
---  -
 0   age             1338 non-null   int64   
 1   sex             1338 non-null   int64   
 2   bmi             1338 non-null   float64  
 3   children        1338 non-null   int64   
 4   smoker          1338 non-null   int64   
 5   region          1338 non-null   int64   
 6   charges         1338 non-null   float64  
 7   insuranceclaim  1338 non-null   int64   
dtypes: float64(2), int64(6)
memory usage: 83.8 KB
```

```
In [6]: df.shape
```

```
Out[6]: (1338, 8)
```


```
In [7]: df.describe()
```

```
Out[7]:
```

	age	sex	bmi	children	smoker	region	charg
count	1338.000000	1338.000000	1338.000000	1338.000000	1338.000000	1338.000000	1338.0000
mean	39.207025	0.505232	30.663397	1.094918	0.204783	1.515695	13270.4222
std	14.049960	0.500160	6.098187	1.205493	0.403694	1.104885	12110.0112
min	18.000000	0.000000	15.960000	0.000000	0.000000	0.000000	1121.8739
25%	27.000000	0.000000	26.296250	0.000000	0.000000	1.000000	4740.2871
50%	39.000000	1.000000	30.400000	1.000000	0.000000	2.000000	9382.0330
75%	51.000000	1.000000	34.693750	2.000000	0.000000	2.000000	16639.9125
max	64.000000	1.000000	53.130000	5.000000	1.000000	3.000000	63770.4280


```
In [8]: sns.boxplot(df["bmi"])
```

```
Out[8]: <Axes: >
```




```
In [9]: sns.countplot(x=df["insuranceclaim"])
```

```
Out[9]: <Axes: xlabel='insuranceclaim', ylabel='count'>
```



```
In [10]: sns.histplot(df["bmi"])
```

```
Out[10]: <Axes: xlabel='bmi', ylabel='Count'>
```



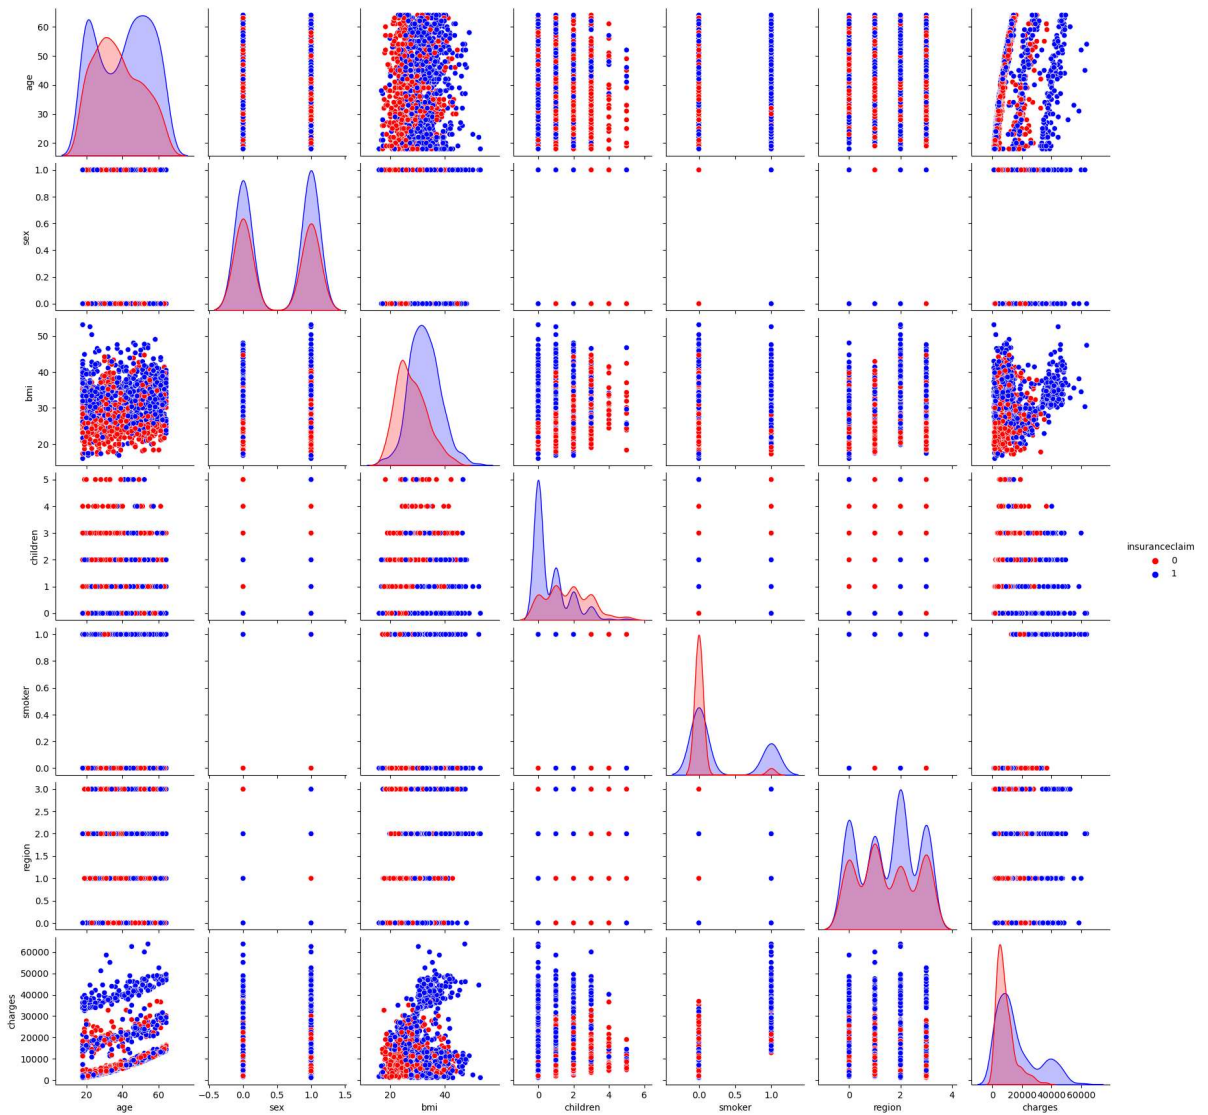
```
In [11]: plt.figure(figsize=(10,10))
sns.heatmap(df.corr(),annot=True,cmap="crest")
```

```
Out[11]: <Axes: >
```



```
In [12]: import warnings
warnings.filterwarnings("ignore")
custom_palette=["Red", "Blue", "Violet"]
sns.pairplot(df, hue="insuranceclaim", palette=custom_palette)
```

Out[12]: <seaborn.axisgrid.PairGrid at 0x16a39135c50>



```
In [13]: q1 = df["bmi"].quantile(0.25) #25 quantile
q2 = df["bmi"].quantile(0.50) #50 quantile
q3 = df["bmi"].quantile(0.75) #75 quantile

#iqr: Inter Quartile Range
iqr = q3-q1
```

```
upper_tail = q3 + 1.5 * iqr
lower_tail = q1 - 1.5 * iqr
```

```
In [14]: upper_tail
```


Out[14]: 47.290000000000006

```
In [15]: lower_tail
```

```
Out[15]: 13.7
```


```
In [16]: sns.kdeplot(df["region"],fill=True)
```

```
Out[16]: <Axes: xlabel='region', ylabel='Density'>
```



```
In [17]: sns.kdeplot(df["children"])
```

```
Out[17]: <Axes: xlabel='children', ylabel='Density'>
```



## Checking Outliers(Function)

```
In [18]: def outliers(n):  
    q1=np.percentile(n,25) #25 Quantile  
    q2=np.percentile(n,50) #50 Quantile  
    q3=np.percentile(n,75) #75 Quantile  
  
    #iqr=Inter Quantile Range  
    iqr=q3-q1  
  
    upper_tail = q3 + 1.5 * iqr  
    lower_tail = q1 - 1.5 * iqr  
  
    outliers = [i for i in n if i < lower_tail or i > upper_tail]  
    return outliers
```

```
In [19]: n=df["bmi"]  
outliers(n)
```

```
Out[19]: [49.06, 48.07, 47.52, 47.41, 50.38, 47.6, 52.58, 47.74, 53.13]
```



```
In [20]: df.corr()
```

```
Out[20]:
```

	age	sex	bmi	children	smoker	region	charges	insuranceclaim
age	1.000000	-0.020856	0.109272	0.042469	-0.025019	0.002127	0.299008	0.113723
sex	-0.020856	1.000000	0.046371	0.017163	0.076185	0.004588	0.057292	0.031565
bmi	0.109272	0.046371	1.000000	0.012759	0.003750	0.157566	0.198341	0.384198
children	0.042469	0.017163	0.012759	1.000000	0.007673	0.016569	0.067998	-0.409526
smoker	-0.025019	0.076185	0.003750	0.007673	1.000000	-0.002181	0.787251	0.333261
region	0.002127	0.004588	0.157566	0.016569	-0.002181	1.000000	-0.006208	0.020891
charges	0.299008	0.057292	0.198341	0.067998	0.787251	-0.006208	1.000000	0.309418
insuranceclaim	0.113723	0.031565	0.384198	-0.409526	0.333261	0.020891	0.309418	1.000000

```
In [27]: df["insuranceclaim"].unique()
```

```
Out[27]: array([1, 0], dtype=int64)
```

## 4: Feature Engineering

### Handling Outliers

```
In [21]: def find_iqr(df,n):
          q1=df[n].quantile(0.25)
          q3=df[n].quantile(0.75)
          iqr=q3-q1
          upper_tail = q3 + 1.5 * iqr
          lower_tail = q1 - 1.5 * iqr
          print(lower_tail)
          print(upper_tail)
          df[n]=df[n].apply(lambda x: upper_tail if x>upper_tail else(lower_tail if x<lower_tail else x),axis=0)
          n="bmi"
          find_iqr(df,n)
```

```
13.7
```

```
47.290000000000006
```

```
In [22]: from sklearn.preprocessing import StandardScaler
std = StandardScaler()
std_array = std.fit_transform(df.iloc[:, :7])
df1 = pd.DataFrame(std_array)
df1
```

```
Out[22]:
```

	0	1	2	3	4	5	6
0	-1.438764	-1.010519	-0.454201	-0.908614	1.970587	1.343905	0.298584
1	-1.509965	0.989591	0.515300	-0.078767	-0.507463	0.438495	-0.953689
2	-0.797954	0.989591	0.388125	1.580926	-0.507463	0.438495	-0.728675
3	-0.441948	0.989591	-1.312218	-0.908614	-0.507463	-0.466915	0.719843
4	-0.513149	0.989591	-0.292342	-0.908614	-0.507463	-0.466915	-0.776802
...	...	...	...	...	...	...	...
1333	0.768473	0.989591	0.052846	1.580926	-0.507463	-0.466915	-0.220551
1334	-1.509965	-1.010519	0.209750	-0.908614	-0.507463	-1.372326	-0.914002
1335	-1.509965	-1.010519	1.023999	-0.908614	-0.507463	0.438495	-0.961596
1336	-1.296362	-1.010519	-0.801041	-0.908614	-0.507463	1.343905	-0.930362
1337	1.551686	-1.010519	-0.260961	-0.908614	1.970587	-0.466915	1.311053

1338 rows × 7 columns

```
In [23]: df.shape
```

```
Out[23]: (1338, 8)
```

```
In [24]: x=df1
y=df["insuranceclaim"]
x
```

```
Out[24]:
```

	0	1	2	3	4	5	6
0	-1.438764	-1.010519	-0.454201	-0.908614	1.970587	1.343905	0.298584
1	-1.509965	0.989591	0.515300	-0.078767	-0.507463	0.438495	-0.953689
2	-0.797954	0.989591	0.388125	1.580926	-0.507463	0.438495	-0.728675
3	-0.441948	0.989591	-1.312218	-0.908614	-0.507463	-0.466915	0.719843
4	-0.513149	0.989591	-0.292342	-0.908614	-0.507463	-0.466915	-0.776802
...	...	...	...	...	...	...	...
1333	0.768473	0.989591	0.052846	1.580926	-0.507463	-0.466915	-0.220551
1334	-1.509965	-1.010519	0.209750	-0.908614	-0.507463	-1.372326	-0.914002
1335	-1.509965	-1.010519	1.023999	-0.908614	-0.507463	0.438495	-0.961596
1336	-1.296362	-1.010519	-0.801041	-0.908614	-0.507463	1.343905	-0.930362
1337	1.551686	-1.010519	-0.260961	-0.908614	1.970587	-0.466915	1.311053

1338 rows × 7 columns

```
In [25]: y
```

```
Out[25]:
```

0	1
1	1
2	0
3	0
4	1
...	..
1333	0
1334	1
1335	1
1336	0
1337	1

Name: insuranceclaim, Length: 1338, dtype: int64

## 5: Model Training

```
In [26]: xtrain,xtest,ytrain,ytest=train_test_split(x,y,test_size=0.2,random_state=10,st
```

```
In [27]: log_clf=LogisticRegression()  
log_clf_model=log_clf.fit(xtrain,ytrain)  
log_clf_model
```

Out[27]: LogisticRegression()

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.**

**On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

## 6: Model Evaluation

```
In [28]: ytrain_predict=log_clf_model.predict(xtrain) #training
```

```
In [29]: accuracy=accuracy_score(ytrain,ytrain_predict)  
print(f"Accuracy : {accuracy}")  
print(""*58)  
  
conf_matrix = confusion_matrix(ytrain,ytrain_predict)  
print(conf_matrix)  
print(""*58)  
  
class_report = classification_report(ytrain,ytrain_predict)  
print(class_report)  
print(""*58)
```

Accuracy : 0.8897196261682243

\*\*\*\*\*

```
[[381  63]  
 [ 55 571]]
```

\*\*\*\*\*

	precision	recall	f1-score	support
0	0.87	0.86	0.87	444
1	0.90	0.91	0.91	626
accuracy			0.89	1070
macro avg	0.89	0.89	0.89	1070
weighted avg	0.89	0.89	0.89	1070

\*\*\*\*\*

```
In [31]: ytest_predict=log_clf_model.predict(xtest) #testing
```

```
In [32]: accuracy=accuracy_score(ytest,ytest_predict)
print(f"Accuracy : {accuracy}")
print(""*58)

conf_matrix = confusion_matrix(ytest,ytest_predict)
print(conf_matrix)
print(""*58)

class_report = classification_report(ytest,ytest_predict)
print(class_report)
print(""*58)
```

Accuracy : 0.8768656716417911

\*\*\*\*\*

```
[[ 91  20]
 [ 13 144]]
*****
```

	precision	recall	f1-score	support
0	0.88	0.82	0.85	111
1	0.88	0.92	0.90	157
accuracy			0.88	268
macro avg	0.88	0.87	0.87	268
weighted avg	0.88	0.88	0.88	268

\*\*\*\*\*

## DECISION TREE

```
In [38]: dt=DecisionTreeClassifier()
dt
```

Out[38]: DecisionTreeClassifier()

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.**

**On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [40]: dt_model= dt.fit(xtrain,ytrain)
dt_model
```

Out[40]: DecisionTreeClassifier()

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.**

**On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [46]: ytrain_predict1 = dt_model.predict(xtrain)  #training
```

```
In [49]: accuracy=accuracy_score(ytrain,ytrain_predict1)
print(f"Accuracy : {accuracy}")
print(""*58)

conf_matrix = confusion_matrix(ytrain,ytrain_predict1)
print(conf_matrix)
print(""*58)

class_report = classification_report(ytrain,ytrain_predict1)
print(class_report)
print(""*58)
```

Accuracy : 1.0

\*\*\*\*\*

```
[[444  0]
 [ 0 626]]
```

\*\*\*\*\*

	precision	recall	f1-score	support
0	1.00	1.00	1.00	444
1	1.00	1.00	1.00	626
accuracy			1.00	1070
macro avg	1.00	1.00	1.00	1070
weighted avg	1.00	1.00	1.00	1070

\*\*\*\*\*

```
In [50]: ytest_predict1=dt_model.predict(xtrain)  #testing
```

```
In [51]: accuracy = accuracy_score(ytest,ytest_predict1)
print(f"Accuracy : {accuracy}")
print(""*55)

conf_matrix = confusion_matrix(ytest,ytest_predict1)
print(conf_matrix)
print(""*55)

class_report = classification_report(ytest,ytest_predict1)
print(class_report)
print(""*55)
```

```
Accuracy : 0.9514925373134329
*****
[[103   8]
 [  5 152]]
*****
              precision    recall  f1-score   support

         0           0.95        0.93        0.94         111
         1           0.95        0.97        0.96         157


 accuracy                   0.95         268
 macro avg           0.95        0.95        0.95         268
 weighted avg        0.95        0.95        0.95         268

*****
```


```
In [52]: dt
```

```
Out[52]: DecisionTreeClassifier()
```

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.**

**On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [54]: dt_plot = plot_tree(dt_model,class_names=["Yes","No"], filled = True)
```



## DECISION TREE PARAMETERS

```
In [55]: hyperparameters={"criterion":["entropy","gini"],
                          "max_depth":np.arange(1,10),
                          "min_samples_split":np.arange(2,10),
                          "min_samples_leaf":np.arange(2,10),}
```

```
In [59]: random_search_model=RandomizedSearchCV(dt_model,hyperparameters,error_score="raise",
random_search_model
```

```
Out[59]: RandomizedSearchCV(cv=5, error_score='raise',
                          estimator=DecisionTreeClassifier(),
                          param_distributions={'criterion': ['entropy', 'gini'],
                                              'max_depth': array([1, 2, 3, 4, 5, 6,
7, 8, 9]),
                                              'min_samples_leaf': array([2, 3, 4,
5, 6, 7, 8, 9]),
                                              'min_samples_split': array([2, 3, 4,
5, 6, 7, 8, 9])})
```

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.**

**On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**



```
In [60]: random_search_model.fit(xtrain,ytrain)
```

```
Out[60]: RandomizedSearchCV(cv=5, error_score='raise',
                             estimator=DecisionTreeClassifier(),
                             param_distributions={'criterion': ['entropy', 'gini'],
                                                  'max_depth': array([1, 2, 3, 4, 5, 6,
                                                                    7, 8, 9]),
                                                  'min_samples_leaf': array([2, 3, 4,
                                                                    5, 6, 7, 8, 9]),
                                                  'min_samples_split': array([2, 3, 4,
                                                                    5, 6, 7, 8, 9])})
```

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.**

**On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [62]: random_search_model.best_estimator_
```

```
Out[62]: DecisionTreeClassifier(criterion='entropy', max_depth=9, min_samples_leaf=5,
                                min_samples_split=5)
```

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.**

**On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [63]: hyp_dt_model = DecisionTreeClassifier(criterion='entropy', max_depth=3, min_sam
```

```
In [64]: hyp_dt_model.fit(xtrain,ytrain)
```

```
Out[64]: DecisionTreeClassifier(criterion='entropy', max_depth=3, min_samples_leaf=13)
```

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.**

**On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [73]: ypred_train = hyp_dt_model.predict(xtrain)
```


```
In [74]: print(f"Accuracy = {accuracy_score(ytrain,ypred_train)}") #TESTING
print(classification_report(ytrain,ypred_train))
print(sns.heatmap(confusion_matrix(ytrain,ypred_train),annot=True))
```

```
Accuracy = 0.8429906542056075
      precision    recall  f1-score   support

     0       0.73      0.99      0.84       444
     1       0.99      0.74      0.85       626

 accuracy          0.84       1070
 macro avg       0.86      0.86      0.84       1070
 weighted avg    0.88      0.84      0.84       1070
```

```
Axes(0.125,0.11;0.62x0.77)
```



```
In [75]: ypred_test = hyp_dt_model.predict(xtest)
```

```
In [76]: print(f"Accuracy = {accuracy_score(ytest,ypred_test)}")
print(confusion_matrix(ytest,ypred_test))
print(classification_report(ypred_test,ytest))
```

Accuracy = 0.8432835820895522

[[110 1]

[ 41 116]]

	precision	recall	f1-score	support
0	0.99	0.73	0.84	151
1	0.74	0.99	0.85	117
accuracy			0.84	268
macro avg	0.86	0.86	0.84	268
weighted avg	0.88	0.84	0.84	268

## RANDOM FOREST

```
In [77]: rf_clf = RandomForestClassifier()
```

```
In [80]: rf_clf_model = rf_clf.fit(xtrain,ytrain)
```


```
In [83]: ypred_train1 = rf_clf_model.predict(xtrain)
```

```
In [84]: print("Accuracy = ",accuracy_score(ytrain,ypred_train1)) #TRAINING
sns.heatmap(confusion_matrix(ytrain,ypred_train1),annot=True)
print(classification_report(ytrain,ypred_train1))
```

```
Accuracy = 1.0
          precision    recall  f1-score   support

     0       1.00      1.00      1.00       444
     1       1.00      1.00      1.00       626

 accuracy          1.00          1.00          1.00       1070
 macro avg          1.00          1.00          1.00       1070
weighted avg          1.00          1.00          1.00       1070
```




```
In [85]: ypred_test1 = rf_clf_model.predict(xtest)
```

```
In [86]: print(f"Accuracy = {accuracy_score(ytest,ypred_test1)}") #TESTING
sns.heatmap(confusion_matrix(ytest,ypred_test1),annot=True)
print(classification_report(ytest,ypred_test1))
```

```
Accuracy = 0.9589552238805971
              precision    recall  f1-score   support

      0       0.96       0.94       0.95        111
      1       0.96       0.97       0.97        157

 accuracy          0.96          0.96          0.96          268
 macro avg         0.96          0.96          0.96          268
 weighted avg      0.96          0.96          0.96          268
```



## RANDOM FOREST PARAMETERS

```
In [87]: hyp={
    "n_estimators":np.arange(10,200),
    "criterion": ['gini','entropy'],
    "max_depth" : np.arange(5,15),
    "min_samples_split":np.arange(5,20),
    "min_samples_leaf":np.arange(4,15),
    "max_features":['auto']
}
```

```
In [88]: rscv_rf_clf = RandomizedSearchCV(rf_clf_model, hyp, cv = 6)
```

```
In [91]: rscv_rf_clf.fit(xtrain,ytrain)
```

```
Out[91]: RandomizedSearchCV(cv=6, estimator=RandomForestClassifier(),
                             param_distributions={'criterion': ['gini', 'entropy'],
                                                  'max_depth': array([ 5,  6,  7,  8,
9, 10, 11, 12, 13, 14]),
                                                  'max_features': ['auto'],
                                                  'min_samples_leaf': array([ 4,  5,
6,  7,  8,  9, 10, 11, 12, 13, 14]),
                                                  'min_samples_split': array([ 5,  6,
7,  8,  9, 10, 11, 12, 13, 14, 15, 16, 17, 18, 19]),
                                                  'n_estimators': array([ 10,  11,  12,
13,  14,  15,  16,  17...
101, 102, 103, 104, 105, 106, 107, 108, 109, 110, 111, 112, 113,
114, 115, 116, 117, 118, 119, 120, 121, 122, 123, 124, 125, 126,
127, 128, 129, 130, 131, 132, 133, 134, 135, 136, 137, 138, 139,
140, 141, 142, 143, 144, 145, 146, 147, 148, 149, 150, 151, 152,
153, 154, 155, 156, 157, 158, 159, 160, 161, 162, 163, 164, 165,
166, 167, 168, 169, 170, 171, 172, 173, 174, 175, 176, 177, 178,
179, 180, 181, 182, 183, 184, 185, 186, 187, 188, 189, 190, 191,
192, 193, 194, 195, 196, 197, 198, 199])})
```

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.**

**On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [92]: rscv_rf_clf.best_estimator_
```

```
Out[92]: RandomForestClassifier(criterion='entropy', max_depth=14, max_features='auto',
                                min_samples_leaf=4, min_samples_split=12,
                                n_estimators=128)
```

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.**

**On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [93]: rf_clf = RandomForestClassifier(criterion='entropy', max_depth=13, max_features=
        min_samples_leaf=10, min_samples_split=7,
        n_estimators=135, random_state=10)

rf_clf
```

```
Out[93]: RandomForestClassifier(criterion='entropy', max_depth=13, max_features='auto',
                                min_samples_leaf=10, min_samples_split=7,
                                n_estimators=135, random_state=10)
```

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.**

**On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [94]: rf_clf.fit(xtrain,ytrain)
```

```
Out[94]: RandomForestClassifier(criterion='entropy', max_depth=13, max_features='auto',
                                min_samples_leaf=10, min_samples_split=7,
                                n_estimators=135, random_state=10)
```

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.**

**On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [95]: y_pred_train = rf_clf.predict(xtrain)
```

```
In [96]: accuracy = accuracy_score(ytrain,y_pred_train)                                #TRAINING
print(f"Accuracy : {accuracy}")
conf_matrix = confusion_matrix(ytrain,y_pred_train)
print(conf_matrix)
class_report = classification_report(ytrain,y_pred_train)
print(class_report)
```

Accuracy : 0.9672897196261683

[[428 16]

[ 19 607]]

	precision	recall	f1-score	support
0	0.96	0.96	0.96	444
1	0.97	0.97	0.97	626
accuracy			0.97	1070
macro avg	0.97	0.97	0.97	1070
weighted avg	0.97	0.97	0.97	1070

```
In [97]: y_pred_test = rf_clf.predict(xtest)
```

```
In [98]: accuracy = accuracy_score(ytest,y_pred_test)           #TESTING
print(f"Accuracy : {accuracy}")
conf_matrix = confusion_matrix(ytest,y_pred_test)
print(conf_matrix)
class_report = classification_report(ytest,y_pred_test)
print(class_report)
```

```
Accuracy : 0.9402985074626866
[[101  10]
 [  6 151]]
```

	precision	recall	f1-score	support
0	0.94	0.91	0.93	111
1	0.94	0.96	0.95	157
accuracy			0.94	268
macro avg	0.94	0.94	0.94	268
weighted avg	0.94	0.94	0.94	268

## ADABOOST

```
In [103]: adb_clf = AdaBoostClassifier()
adb_clf
```

Out[103]: AdaBoostClassifier()

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.**

**On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [104]: adb_clf_model = adb_clf.fit(xtrain,ytrain)
adb_clf_model
```

Out[104]: AdaBoostClassifier()

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.**

**On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [105]: ypred__train1 = adb_clf_model.predict(xtrain)
```



```
In [108]: accuracy = accuracy_score(ytrain,ypred__train1)           #TRAINING
print(f"Accuracy : {accuracy}")

conf_matrix = confusion_matrix(ytrain,ypred__train1)
print(conf_matrix)

class_report = classification_report(ytrain,ypred__train1)
print(class_report)
```

```
Accuracy : 0.8897196261682243
[[384  60]
 [ 58 568]]

      precision    recall  f1-score   support

     0       0.87       0.86       0.87         444
     1       0.90       0.91       0.91         626

 accuracy          0.89          1070
 macro avg         0.89          1070
 weighted avg      0.89          1070
```

```
In [107]: ypred__test1 = adb_clf_model.predict(xtest)
```

```
In [109]: accuracy = accuracy_score(ytest,ypred__test1)           #TRAINING
print(f"Accuracy : {accuracy}")

conf_matrix = confusion_matrix(ytest,ypred__test1)
print(conf_matrix)

class_report = classification_report(ytest,ypred__test1)
print(class_report)
```

```
Accuracy : 0.8283582089552238
[[ 84  27]
 [ 19 138]]

      precision    recall  f1-score   support

     0       0.82       0.76       0.79         111
     1       0.84       0.88       0.86         157

 accuracy          0.83          268
 macro avg         0.83          268
 weighted avg      0.83          268
```

## ADABOOST HYPERPARAMETERS

```
In [112]: hyp = {"n_estimators":np.arange(10,200),  
               "max_depth": [None, 10, 20, 30]}  
  
rscv_adb_clf_model = RandomizedSearchCV(rf_clf_model, hyp, cv = 6)  
rscv_adb_clf_model.fit(xtrain,ytrain)  
rscv_adb_clf_model.best_estimator_
```

Out[112]: RandomForestClassifier(max\_depth=30, n\_estimators=51)

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.**

**On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [120]: adb_clf = AdaBoostClassifier(learning_rate=0.1, n_estimators=36, random_state=10)  
adb_clf
```

Out[120]: AdaBoostClassifier(learning\_rate=0.1, n\_estimators=36, random\_state=10)

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.**

**On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [121]: adb_clf.fit(xtrain,ytrain)
```

Out[121]: AdaBoostClassifier(learning\_rate=0.1, n\_estimators=36, random\_state=10)

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.**

**On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [122]: y_pred__train = adb_clf.predict(xtrain)
```

```
In [124]: accuracy = accuracy_score(ytrain,y_pred__train)           #TRAINING
print(f"Accuracy : {accuracy}")

conf_matrix = confusion_matrix(ytrain,y_pred__train)
print(conf_matrix)

class_report = classification_report(ytrain,y_pred__train)
print(class_report)
```

```
Accuracy : 0.8411214953271028
[[334 110]
 [ 60 566]]
```

	precision	recall	f1-score	support
0	0.85	0.75	0.80	444
1	0.84	0.90	0.87	626
accuracy			0.84	1070
macro avg	0.84	0.83	0.83	1070
weighted avg	0.84	0.84	0.84	1070

```
In [125]: y_pred__test = adb_clf.predict(xtest)
```

```
In [128]: accuracy = accuracy_score(ytest,y_pred__test)           #TESTING
print(f"Accuracy : {accuracy}")

conf_matrix = confusion_matrix(ytest,y_pred__test)
print(conf_matrix)

class_report = classification_report(ytest,y_pred__test)
print(class_report)
```

```
Accuracy : 0.7910447761194029
[[ 72  39]
 [ 17 140]]
```

	precision	recall	f1-score	support
0	0.81	0.65	0.72	111
1	0.78	0.89	0.83	157
accuracy			0.79	268
macro avg	0.80	0.77	0.78	268
weighted avg	0.79	0.79	0.79	268

```
In [ ]:
```